

Georgy Savva

georgy.savva@nyu.edu | georgysavva.github.io | github.com/georgysavva

RESEARCH INTERESTS

World Models, Representation Learning, Digital Agents, Multimodal Learning

EDUCATION

New York University <i>Master of Science in Computer Science, Courant Institute of Mathematical Sciences.</i> GPA: 3.8	New York, United States 2024 – Expected 2025
Russian University of Cooperation <i>Bachelor of Science in Computer Science</i>	Moscow, Russia 2019 – 2021

RESEARCH EXPERIENCE

Research Assistant <i>NYU Courant Institute of Mathematical Sciences, Advised by Prof. Saining Xie</i>	Jan 2025 – Present New York, United States
• Developed a Minecraft World Model in JAX, matching the performance of the proprietary SOTA model • Developed the first multi-agent Minecraft World Model	
Research Intern <i>NYU Courant Institute of Mathematical Sciences, Advised by Prof. Lerrel Pinto</i>	Jun 2024 – Nov 2024 New York, United States
• Applied online reinforcement learning to fix the human demonstration morphology gap for dexterous robotic manipulation • Trained a behavior cloning policy using transformers on third-person human demonstration data to solve dexterous robotic hand tasks • Setup an auto-resettable reinforcement learning environment for online robot training using the spacemouse device	

TEACHING EXPERIENCE

Graduate Teaching Assistant <i>NYU Courant Institute of Mathematical Sciences. Prof. Saining Xie</i>	Sep 2025 – Dec 2025 New York, United States
• Led weekly office hours for CSCI-GA.2271-001 Computer Vision. Website • Helped students with their research projects and homework	

PUBLICATIONS

1. *Solaris: A Multi-agent Video World Model*
G. Savva*, O. Michel*, P. Waiwitlikhit*, D. Lu*, T. Meehan, D. Mishra, J. Lu, S. Poddar, S. Xie.
(Work in Progress).
2. *HuDOR: Bridging the Human to Robot Dexterity Gap through Object-Oriented Rewards*
I. Guzey, Y. Dai, **G. Savva**, R. Bhirangi, L. Pinto.
ICRA 2025 (Accepted). [Website](#)

PROJECTS

World Model Finetuning Analysis Investigated the generalization and finetuning of two SOTA Minecraft World Models on three data distributions. Website	Jan 2025 – May 2025
AppSim: A Learned World Model for an App API Used ChatGPT o3 in a zero-shot setting, achieving 74% accuracy, and compared its performance to a finetuned TinyLlama model. Website	Jan 2025 – May 2025
Transformer-Based Diffusion for Game Generation Trained a transformer diffusion model to simulate DOOM trajectories. It achieves a PSNR of 32.21 in the teacher-forcing setting, producing an indistinguishable quality from the ground truth. Website	Oct 2024 – Dec 2024

Advantage Actor-Critic with Optuna	Jul 2024 – Sep 2024
Trained a reinforcement learning agent with Optuna for the HalfCheetah env, achieving a 24% better performance than the best publicly available policy. Website	
Open-Source Library Scany	May 2020 – Present
Created an open-source library to map data from a database into Go objects. The library has 1,300 stars on GitHub and is used by thousands of companies. GitHub	

INDUSTRY EXPERIENCE

Senior Software Engineer	Oct 2023 – May 2024
<i>Raylu</i>	<i>New York, United States</i>
<ul style="list-style-type: none"> Led backend development of an LLM chatbot for healthcare and deployed it to production, serving first paying customers Developed an LLM-powered SaaS workflow automation application using the open source workflow engine activepieces, allowing the company to release a new product in under 3 months 	
Technical Co-Founder	Jul 2022 – Aug 2023
<i>Scifind</i>	<i>Los Angeles, United States</i>
<ul style="list-style-type: none"> Launched a troubleshooting platform for bioscientists into production, gaining 4,000 MAU in the first 2 months Led the development of the product with a team of 2 engineers using TypeScript, Next.js, and Node.js 	
Senior Software Engineer	Jan 2021 – Jul 2022
<i>IOTA</i>	<i>Berlin, Germany</i>
<ul style="list-style-type: none"> Developed a new network layer of the blockchain node using libp2p, decreasing the number of peering errors by 70% Set up automatic deployment via Ansible and GitHub Actions, accelerating the development speed by 25% Introduced Go guidelines to the project and tools to ensure them, reducing the frequency of bugs by 50% 	
Senior Software Engineer	Jan 2020 – Jan 2021
<i>Elsa</i>	<i>San Francisco, United States</i>
<ul style="list-style-type: none"> Implemented App Store and Google Play in-app subscriptions as microservices using FastAPI, which improved the correctness of billing by 50% Integrated payment gateway Instamojo, increasing the number of app purchases in Asia threefold 	
Software Engineer	Mar 2017 – May 2019
<i>Edwin.ai</i>	<i>San Francisco, United States</i>
<ul style="list-style-type: none"> Participated in building the English tutor chatbot from a prototype stage to being the top 1 bot on Facebook Implemented a YAML-file-based framework for building dialog systems, which allowed the company to launch a new product on Google Assistant Designed a distributed queue using PostgreSQL to handle messages from 800,000 users Developed the user knowledge graph service using Neo4j, improving the performance of read queries threefold 	
Software Engineer	Jun 2016 – Mar 2017
<i>Snaappy</i>	<i>Moscow, Russia</i>
<ul style="list-style-type: none"> Rewrote the messaging service in Go, which increased the backend performance fivefold Migrated the chat storage from PostgreSQL to MongoDB, allowing horizontal scalability for the data layer Implemented a Websocket service for real-time updates using Go, driving user engagement by 30% 	

SKILLS

Languages: Python, SQL, Go, TypeScript
Machine Learning: Transformers, Diffusion, U-Net, CNN, LSTM, Reinforcement Learning, Distributed Training, TPU
Libraries: PyTorch, JAX, Gym, Optuna, Hydra, TensorFlow, Numpy, Pandas, FastAPI, React, Node.js
Tools: Docker, Kubernetes, PostgreSQL, AWS, GraphQL, Hasura

CERTIFICATIONS

Probability & Statistics for Machine Learning & Data Science	Sep 2024
<i>DeepLearning.AI, GPA: 93/100</i>	
YC Summer Startup School	Aug 2022
<i>Y Combinator</i>	
Mathematics for Machine Learning: Multivariate Calculus	Dec 2021
<i>Imperial College London, GPA: 96/100</i>	
Mathematics for Machine Learning: Linear Algebra	Dec 2021
<i>Imperial College London, GPA: 96/100</i>	